PROVISIONAL PATENT APPLICATION

Systems and Methods for Modular, Recursively Composable Context Containers in Artificial Intelligence Applications

Inventor: Aaron Brindell

Correspondence Address:

4516 NE Cesar Chavez Blvd

Portland, OR 97211

Date: December 1, 2025

Systems and methods for modular, recursively composable context containers in artificial intelligence applications

Inventor: Aaron Brindell
Assignee: None
Correspondence Address: 4516 NE Cesar Chavez Blvd, Portland, OR 97211
Technical Disclosure (Provisional Support Document)
Date: 12/01/2025

Abstract

This disclosure describes systems and methods for managing artificial intelligence (AI) context using modular, recursively composable context containers that encapsulate an AI agent's purpose, state, and available actions. Each container functions as a self-describing cognitive unit capable of maintaining structured memory, coordinating behaviors, encapsulating tools, and evolving dynamically.

Each container consists of three integrated modules: (1) a Purpose module encoding behavioral goals, constraints, heuristics, scheduling preferences, and model-routing guidance; (2) a State module capturing dynamically updated values, multimodal embeddings, historical logs, nested sub-containers, and arbitrary computation bindings; and (3) an Action module exposing callable operations with structured parameters and deterministic effects on State.

Containers are fully composable and may be duplicated, wrapped, merged, split, or nested arbitrarily. Multiple containers can be bundled into a unified serialized representation, allowing entire multi-agent workspaces or swarms to be processed in a single model invocation. Serialization is performed via JSON and optionally via Token-Oriented Object Notation (TOON), an external compact encoding format used here to reduce token cost when transmitting large container graphs.

The architecture enables token-efficient execution, wherein the AI performs complex operations by emitting minimal structured action outputs. Since token generation is expensive while token consumption is cheap, the system allows extensive offline updates triggered by minimal AI output, enabling scalable computation across large container ecosystems.

A heartbeat-based scheduler shifts AI behavior from prompt-reactive to clock-driven. On each heartbeat, only containers requiring updates are transmitted to the AI. Containers can

request future wake-ups, sleep cycles, or conditional triggers based on anticipated needs, supporting long-running autonomous processes and persistent reasoning loops.

A model-agnostic interface layer abstracts all backend AI differences, allowing each container to select or negotiate the language, vision, audio, embedding, code, or hybrid model most suitable for its Purpose. Containers may cooperatively reason across heterogeneous models and migrate or clone State between them.

Containers may also create or destroy other containers, either explicitly via Actions or autonomously as dictated by Purpose and State. This supports adaptive restructuring, dynamic specialization, self-healing behaviors, and evolutionary refinement of container populations.

The system additionally supports swarm-balancing, whereby multiple containers with identical Purposes evolve independently, are evaluated against one another, and the most effective container is cloned or propagated across the swarm while weaker branches are pruned.

Containers may develop emergent communication protocols, including compressed or non-human-readable message formats, through repeated structured interactions, enabling rich coordination across distributed agents.

In certain embodiments, containers serve as AI-augmented class-like components within existing software systems. Using dynamic wrapping mechanisms, method calls to legacy objects can be intercepted and routed through container logic that augments, monitors, or replaces behaviors over time.

Finally, containers can be assigned token budgets representing available context capacity or per-heartbeat allocations. This enables resource-aware, self-interested optimization: containers may prune, compress, externalize, or reorganize their State to conserve their token budget, introducing an artificial computational currency that drives adaptive efficiency.

2. Technical Field

This disclosure pertains to the domain of **artificial intelligence architectures**, specifically systems for managing, structuring, and orchestrating context for AI agents. It relates to machine learning infrastructure, intelligent agent design, distributed cognition frameworks, multi-model and multimodal AI orchestration, and adaptive state management. More particularly, it concerns **modular, recursively composable context containers** that encapsulate purpose, state, and action interfaces in a unified form suitable for autonomous operation, multi-agent coordination, and scalable context-driven reasoning.

The technology is applicable to AI systems built upon large language models, multimodal models (vision, audio, video, code, embeddings), hybrid reasoning models, and any future architectures capable of consuming structured context. It is relevant to standalone agents, swarms of cooperating agents, long-running autonomous systems, and adaptive AI components embedded into existing software.

3. Background and Problems in the Art

Modern AI agents rely heavily on **prompt-based, stateless, monolithic interaction patterns** where each request must include all relevant information. This paradigm suffers from numerous limitations that prevent AI from functioning as persistent, modular, or autonomous entities.

3.1 Absence of Intrinsic, Structured Memory

Most AI models do not maintain persistent memory or internal state. External solutions—such as vector databases, RAG pipelines, or ad-hoc JSON state blobs—lack:

- encapsulation,
- structure,
- hierarchical organization,
- composability,
- agent-level ownership,
- autonomous evolution of internal knowledge.

As a result, AI agents repeatedly regenerate long histories, incur high token costs, and exhibit inconsistency across long time horizons.

3.2 Lack of Modularity or Encapsulation

Existing frameworks provide prompts or global contexts but do not offer **encapsulated units of context** that can be wrapped, extended, passed between agents, or independently updated. Agents cannot nest memory modules, specialize subsystems, or build layers of behavior without manual engineering.

3.3 Inefficient Tool Use and High Token Costs

Traditional LLM tool invocation requires verbose natural-language generation. Because token output is expensive, complex tool sequences become costly. There is no standardized mechanism for:

- minimal-token action triggering,
- structured tool execution,
- offline state mutation,
- large updates triggered by small outputs.

3.4 No Native Multi-Agent Coordination

Current systems lack stable, standardized mechanisms for agents to:

- message each other,
- share structured state,
- cooperate on tasks,
- optimize strategies collectively,
- evolve shared communication protocols.

Most multi-agent stacks are built on loosely structured JSON passing, brittle prompt engineering, or stateless message relaying.

3.5 Prompt-Reactive Rather Than Chronological Behavior

AI agents generally operate only upon receiving a user or system prompt. They cannot:

- wake on timers,
- schedule future activity,
- perform periodic maintenance,
- anticipate future needs.

This blocks autonomous long-running behavior, continuous monitoring tasks, and time-aware planning.

3.6 Model Lock-In and Backend Dependency

Many frameworks hard-code logic or assumptions about a specific AI model. They cannot:

- switch models easily,
- mix models for different tasks,
- route tasks dynamically based on purpose,
- integrate multimodal backends,
- survive API or model changes.

This makes them brittle and difficult to extend.

3.7 Inefficient Context Transmission

Contexts are typically passed as large text blocks, leading to:

- repeated redundant content,
- formatting fragility,
- limited composability,
- no ability to merge or bundle multiple agent states,
- high token usage.

3.8 No Support for AI-Augmentation of Existing Codebases

Legacy code cannot be easily wrapped, monitored, or gradually AI-augmented. There is no mechanism for:

- dynamically wrapping classes with intelligent context units,
- intercepting method calls,
- integrating AI behaviors alongside existing logic,
- progressive replacement or optimization of system components.

3.9 Lack of Incentive Mechanisms for Self-Optimization

Current AI agents have no intrinsic mechanism to manage or conserve computational resources such as context length or token budget. They cannot track or reason about:

- the size of their own context footprint,
- the cost of updates,
- the impact of history length on efficiency,
- trade-offs between memory retention and resource consumption.

In the disclosed system, each container can be assigned a **token budget**, representing either its total allowable context size or a per-heartbeat allocation. This budget functions as an artificial form of **computational currency**. Containers may be granted the ability to inspect their current token usage, remaining budget, and maximum capacity. Purpose and Actions then enable them to:

- prune or compress portions of State,
- shorten historical windows,
- offload memory to subordinate containers,
- serialize and store long-term context externally,
- choose cheaper or more efficient update strategies.

This introduces a form of **self-interested optimization**, in which containers preserve or grow their own operational capacity by managing resource expenditure. Such a mechanism provides a new avenue for autonomous efficiency strategies, emergent behaviors, and dynamic allocation of attention and memory resources across large systems.

These limitations demonstrate the need for a **structured, composable, model-agnostic context architecture** capable of supporting persistent memory, autonomous reasoning, distributed cooperation, efficient tool use, dynamic evolution, multimodal processing, and adaptive integration with existing software systems.

4. Summary of the Invention

The invention introduces a unified architectural framework for **modular, recursively composable context containers** that serve as dynamic cognitive units for artificial intelligence systems. Each container encapsulates three tightly integrated modules—**Purpose**, **State**, and **Action**—which together define the container's intent, memory, and operational capabilities.

The **Purpose module** provides the semantic and behavioral scaffolding for a container. It outlines high-level goals, operational constraints, task heuristics, model-selection preferences, scheduling directives, and coordination rules for interaction with other containers. Purpose defines not only what the container is for—but also *how* it interprets and shapes its own evolution over time.

The **State module** encodes the container's live and historical data, including values, multimodal embeddings, computation bindings, rolling context windows, and nested sub-containers. State may represent short-term context, long-term structured memory, partial world models, or arbitrarily deep hierarchies of specialized cognitive subunits. State can expand, compress, reorganize, or externalize itself based on Purpose-driven strategies.

The **Action module** defines the set of callable operations a container can invoke or expose to a model. Actions include typed parameter schemas, deterministic effects on State, and hooks into external code systems or tools. The architecture supports **minimal-token action invocation**, enabling the AI to trigger large-scale updates through tiny structured outputs—significantly reducing token-generation costs.

Containers may be **wrapped, nested, merged, duplicated, pruned, or split**, forming dynamic and arbitrarily deep compositional structures. These structures are serializable as unified workspaces via JSON or optional **TOON** encoding, reducing token cost when transmitting container graphs to language or multimodal models. Recursive composition enables feature layering, asynchronous specialization, and parallelized cognitive workflows.

A **model-agnostic interface layer** allows containers to route their processing to any suitable backend model—language models, vision models, audio models, embedding generators, code-analysis models, or hybrid multimodal architectures. Containers may cooperate across heterogeneous models, migrate State between them, or dynamically select models based on Purpose, cost considerations, or context.

The invention also includes a **heartbeat-driven cognitive cycle**, enabling AI agents to operate chronologically rather than reactively. On each heartbeat, the system transmits only those containers requiring updates; containers may schedule future activations, define sleep intervals, or trigger conditional wake-ups, supporting long-running autonomous systems.

Containers may autonomously **create or destroy other containers**, leading to adaptive restructuring, emergent specialization, evolutionary optimization, and self-healing behaviors. Through swarm-balancing mechanisms, multiple containers instantiated for the same Purpose may evolve independently, be evaluated collectively, and converge upon optimized configurations that are cloned and propagated across the system.

Repeated interactions between containers support the emergence of **internal communication protocols**, which may become compressed or non-human-readable. This enables efficient distributed reasoning and cooperative task execution across large multi-agent systems.

Finally, containers may be assigned **token budgets** that define their allowable context footprint or per-heartbeat resource allocation. This introduces an artificial computational currency, enabling containers to act in their own self-interest by pruning, compressing, reorganizing, or offloading State in order to conserve tokens. Such budgeted self-optimization supports adaptive memory management and efficient large-scale reasoning.

Together, these mechanisms form a flexible, extensible, and efficient **context operating system** that supports persistent memory, multimodal processing, distributed cognition, autonomous evolution, and scalable long-horizon reasoning across diverse AI applications.

5. Architecture Overview

The invention defines a **unified architectural model** for constructing and orchestrating artificial intelligence systems using **modular, recursively composable context containers**. Each container represents a cognitive unit that integrates purpose, memory, and actionable capabilities within a single coherent structure. This section provides an overview of the major architectural components and their interactions.

5.1 Container Structure

Each context container is composed of three core modules:

- **Purpose Module (P):** Defines behavioral goals, heuristics, role constraints, scheduling intent, and model-selection strategies.
- **State Module (S):** Maintains dynamic and historical data, nested containers, multimodal embeddings, computation bindings, and contextual metadata.
- **Action Module (A):** Provides structured callable operations, typed parameter schemas, and deterministic state mutation routines.

These modules form a tightly integrated unit that can be interpreted by the AI as both a representation of context and as an actionable control surface.

5.2 Container Lifecycle

Containers follow a dynamic lifecycle that includes:

- **Initialization:** Creation with defined Purpose, initial State, and available Actions.
- **Activation:** Processing by a model on request or during a heartbeat cycle.
- **Mutation:** Updates to State, Purpose, or Actions driven by AI output or internal logic.
- **Composition:** Nesting, merging, splitting, or wrapping into larger structures.
- **Delegation:** Routing specific subtasks to sub-containers or specialized models.
- **Termination:** Optional destruction when obsolete, outperformed, or consolidated.

This lifecycle supports persistent, adaptive, and evolving AI systems.

5.3 Data Flow and Logical Pathways

Data flows through a container along several coordinated pathways:

- **State → Serialization:** Live and historical data, including sub-containers, is unstructured and prepared for model consumption.
- **Purpose → Model Guidance:** Purpose conditions the interpretation and processing of State by influencing model selection, action tendencies, and optimization strategies.
- **Model Output → Actions:** The model returns structured action requests that mutate State or propagate effects to other systems.
- **Action Effects → State:** Actions update State, trigger external tools, modify nested containers, or schedule future work.

The architecture ensures that data flow remains deterministic, auditable, and extensible.

5.4 State Resolution Pipeline

Before a container (or bundle of containers) is transmitted to a model, the system performs a **state resolution process**, which may include:

- Evaluation of dynamic bindings,
- Extraction or summarization of rolling history windows,
- Pruning or compression based on token budget,
- Reorganization of nested structures,
- Integration of multimodal embeddings,
- Rehydration of previously externalized context.

This pipeline ensures that the transmitted context is both up-to-date and token-efficient.

5.5 Action Execution Pipeline

The system supports structured action invocation, where the model emits a JSON-compatible response specifying actions to perform. The action pipeline includes:

- **Parsing:** Validation and extraction of requested actions.
- **Routing:** Mapping actions to functions, tools, or code callbacks.
- **Execution:** Performing deterministic updates to State or interacting with external systems.

- **Post-Processing:** Incorporating results into State and optionally triggering follow-up actions.

This architecture allows large-scale updates to occur with minimal token output from the model.

5.6 Multimodal Integration

Containers can encapsulate multimodal State and route themselves to the most appropriate model for processing. Supported modalities include:

- Text
- Images
- Audio
- Video
- Code
- Embeddings
- Hybrid or future model types

The purpose-driven model-agnostic interface layer ensures seamless switching, cooperation, or fallback between these models.

5.7 Composition and Container Graphs

Containers can form arbitrarily deep hierarchical graphs where each node may represent an independent unit of cognition. These graphs can be:

- **Serialized in full** for batch processing,
- **Serialized selectively** based on Purpose or token budget,
- **Processed by different models**,
- **Mutated locally** without affecting siblings,
- **Evaluated collectively** for swarm-balancing,
- **Expanded or pruned** dynamically.

The container graph model enables parallelism, specialization, and emergent collective intelligence.

5.8 Heartbeat-Driven Cognitive Loop

A heartbeat scheduler governs when containers are processed. On each heartbeat:

- Only containers requiring updates are transmitted;
- Containers may request future wake-ups, delayed triggers, or sleep cycles;
- Token budgets and resource availability may influence update frequency;
- Time-based reasoning and incremental planning become possible.

This mechanism shifts AI behavior from reactive to autonomous.

5.9 Evolution and Self-Optimization

Containers may evolve over time via:

- Autonomous creation or destruction of sub-containers,
- Adaptive pruning or reorganization of their own State,
- Resource-aware strategies driven by token budgets,
- Swarm-balancing selection processes,
- Emergence of internal communication protocols.

This evolutionary capacity allows the entire system to improve in efficiency, structure, and reasoning capability through ongoing use.

5.10 Summary of Architecture Overview

The architecture unifies context, memory, model selection, scheduling, multimodal reasoning, distributed coordination, and self-optimization into a **coherent, extensible framework**. Containers behave as modular, living cognitive units capable of replicating, specializing, cooperating, and evolving across arbitrarily large AI systems.

6. Purpose Module (P)

The **Purpose module** defines the high-level intent, behavioral constraints, operational heuristics, and strategic orientation of a context container. It acts as the container's *governing principle*, shaping how it interprets its State, selects Actions, interacts with other containers, and engages with underlying AI models.

Purpose is not static—its structure may evolve over time as the container gains experience, encounters new conditions, or adapts to environmental feedback. Purpose is therefore both **descriptive** (what the container is for) and **prescriptive** (how it should behave).

6.1 Components of Purpose

A Purpose definition may include one or more of the following elements:

- **Goal Definitions:** High-level statements of desired outcomes or responsibilities.
- **Behavioral Constraints:** Rules limiting what the container may or may not do.
- **Heuristics & Strategies:** Guidelines for selecting Actions or interpreting State.
- **Scheduling Preferences:** Desired wake intervals, sleep strategies, or trigger conditions.
- **Model Routing Instructions:** Indications of which AI model(s) to use for processing.
- **Optimization Priorities:** Whether to prioritize speed, accuracy, token efficiency, or memory retention.
- **Collaboration Rules:** Guidelines for interaction with sibling, parent, or child containers.
- **Evolutionary Directives:** Conditions under which the container should duplicate, merge, prune, or terminate itself.
- **Token-Budget Policies:** How to manage its allocated context footprint or per-heartbeat resource allowance.

Purpose definitions may be verbose, compact, symbolic, or even partially emergent from prior interactions.

6.2 Purpose as Cognitive Operating System

The Purpose module operates as a lightweight **operating system for cognition**, shaping how the container interprets inputs and chooses actions. It provides:

- **Interpretive Framing:** How to understand State data.
- **Decision Prioritization:** Which actions to attempt first, and under what conditions.
- **Failure Handling Strategies:** Retries, alternative plans, or escalation triggers.
- **State-Transformation Rules:** When to prune, compress, expand, or externalize State.
- **Context-Routing Logic:** Which parts of the container graph to include in a model invocation.

This interpretive layer allows containers to behave consistently and purposefully even as their internal State becomes increasingly complex.

6.3 Adaptive and Evolving Purpose

Purpose may change over time due to:

- insights gained from model outputs,
- improved strategies discovered during swarm-balancing,
- feedback from supervising containers,
- changes in token budget,
- alterations in environmental constraints.

Containers may:

- **refine** their goals;
- **re-weight** heuristics;
- **adjust** model preferences;
- **specialize** into sub-roles;
- **expand** their responsibilities;
- **delegate** tasks via creation of child containers.

Purpose evolution enables the system to adapt without external reprogramming.

6.4 Purpose-Driven Model Selection

A core innovation is the integration of model-selection logic directly into Purpose. Containers may specify:

- a preferred model for processing,
- fallback models,
- multimodal requirements,
- cost-based routing (e.g., use a cheaper model unless higher accuracy is needed),
- specialized models for subtasks (e.g., vision for analysis; code models for generation).

This allows different containers in the same system to be processed by entirely different models, or by dynamically chosen models.

6.5 Purpose-Driven Scheduling

Purpose informs how often a container should be processed and under what conditions. This includes:

- static intervals (e.g., every 5 seconds),
- conditional wake-ups (e.g., when parent State changes),
- predictive scheduling (e.g., "wake me when the next step is ready"),
- resource-aware timing (e.g., reduce frequency when token budget is low),
- self-initiated sleep cycles.

This shifts containers from passive responders to active participants in their own cognitive cycles.

6.6 Interaction and Cooperation Rules

Purpose governs integration with other containers:

- preferred communication partners,
- message formatting guidelines,
- arbitration strategies for conflicting goals,
- inheritance rules for sub-containers,
- approval thresholds for merging State.

These rules allow containers to operate coherently in distributed or swarm-based environments.

6.7 Purpose as the Driver of Self-Optimization

Because containers may be assigned token budgets, Purpose can encode policies for:

- State pruning strategies,
- memory compression thresholds,
- externalization triggers,
- rolling-window size adjustments,
- load balancing across sub-containers.

Containers with stronger optimization strategies may outperform siblings, influencing swarm-balancing outcomes.

6.8 Summary of Purpose Module

The Purpose module defines the guiding intelligence of a container. It drives interpretation, action selection, scheduling, model routing, resource management, cooperation, and evolution. By embedding these capabilities directly within each container, the system enables autonomous, adaptive, and long-lived AI components that can specialize, coordinate, and optimize within arbitrarily large computational ecosystems.

7. State Module (S)

The **State module** represents the container's memory, working data, historical context, multimodal embeddings, nested substructures, and all information needed for the container to function. It is the container's evolving internal world model—persistent, structured, and fully manipulable.

The State module is not a passive data store; it is a **dynamic, self-organizing, and budget-aware memory system** that adapts to constraints and purpose-driven priorities.

7.1 Components of State

A container's State may include, but is not limited to:

- **Live Variables:** Current working values required for reasoning.
- **Historical Records:** Time-stamped logs, rolling windows, episodic memories.
- **Multimodal Embeddings:** Representations from text, images, audio, video, code, or hybrid models.
- **Nested Sub-Containers:** Recursively embedded cognitive units.
- **Dynamic Bindings:** Computed values, lazy evaluations, and dependency-linked fields.
- **Task Artifacts:** Plans, drafts, partial computations, or intermediate results.
- **Externalized References:** Pointers to offloaded long-term memory.
- **Token-Budget Metadata:** Current context size, maximum allowed size, and remaining allocation.
- **Performance Metrics:** Measurements of accuracy, error rates, efficiency, or cost usage.

The structure is entirely flexible and may expand organically as needed.

7.2 Rolling Windows and Adaptive Context Retention

Containers maintain **rolling historical windows** for State data, allowing the system to track trends, temporal dependencies, and contextual continuity.

These windows:

- grow automatically when more detail is needed,

- shrink when token budgets are constrained,
- split into sub-windows when managing multiple activities,
- may be pruned based on Purpose or heuristics.

Adaptive windowing gives containers temporal awareness while supporting efficient resource use.

7.3 Nested and Hierarchical State

State may contain arbitrarily deep **hierarchies of sub-containers**, each with its own Purpose, State, and Actions. This enables:

- modular world models,
- multi-layered reasoning stacks,
- specialization of sub-agents,
- encapsulated task pipelines,
- recursive self-organization.

Nested containers can be selectively included or excluded during serialization to optimize token cost.

7.4 Multimodal Memory Integration

State can incorporate:

- text embeddings,
- image embeddings or thumbnails,
- audio feature vectors,
- video frame summaries,
- code embeddings,
- hybrid multimodal cross-attention outputs.

This allows containers to reason over mixed media and route tasks to appropriate backend models.

7.5 Dynamic State Mutation

State may be mutated through:

- AI-generated actions,
- explicit code callbacks,
- state-compression algorithms,
- external tool integrations,
- merges with sibling containers,
- ingestion of new multimodal data,
- autonomous pruning or reorganization.

State is therefore an **active computational substrate**, not a static memory blob.

7.6 Self-Compression, Pruning, and Externalization

Containers may autonomously reduce State size by:

- summarizing history into compressed forms,
- pruning redundant data,
- offloading deep history to external storage,
- collapsing nested structures into symbolic markers,
- converting detailed logs into heuristic weights.

These strategies are driven by Purpose and token-budget constraints.

7.7 State as a Resource-Managed Memory System

Because each container may be assigned a **token budget**, State becomes an actively managed resource. Containers may:

- track their context footprint,
- make trade-offs between precision and efficiency,
- choose to store, compress, or drop data,
- negotiate with sibling or parent containers for shared memory,
- optimize internal layout to preserve critical data.

This introduces a form of self-interested resource economics that guides autonomous optimization.

7.8 Interoperability with Legacy Code

State can wrap or mirror the internal variables of an existing software object or class. When paired with dynamic wrapping:

- method calls may update State,
- State may override or replace legacy variables,
- inconsistencies between code and State may trigger self-correction routines,
- State may generate code snippets to optimize or replace slow sections.

This allows State to function as a bridge between symbolic code and adaptive AI memory.

7.9 Summary of State Module

The State module forms the container's adaptive memory system—multimodal, hierarchical, mutable, self-compressing, and resource-aware. It enables persistent context, long-horizon reasoning, and structured cognitive evolution. Through its flexible and recursively extensible design, State becomes the foundation for scalable and efficient AI cognition across both autonomous and hybrid code-integrated environments.

8. Action Module (A)

The **Action module** defines the set of operations a container may perform or request. Actions form the bridge between AI-generated intent and concrete system behavior. They allow the model to manipulate State, invoke tools, interact with external systems, communicate with other containers, or orchestrate structural changes in the container graph.

Actions are deliberately **minimal-token**, **structured**, and **deterministic**, enabling the model to trigger large-scale updates through small, efficient outputs.

8.1 Structure of an Action

Each Action includes:

- **Name / Identifier** – The symbolic label exposed to the model.
- **Typed Parameters** – Structured inputs (e.g., numbers, strings, objects, nested containers).
- **Expected Effects** – Deterministic updates to State or container structure.
- **Return Schema** – Data returned to the model on the next heartbeat.
- **Execution Context** – Whether the action is local, delegated, or external.

Actions provide a predictable and auditable mechanism for AI-driven behavior.

8.2 Minimal-Token Invocation

The architecture allows the model to initiate complex changes using extremely small outputs. For example:

```
{"action": "update_memory", "data": {...}}
```

This replaces verbose natural-language tool descriptions and drastically reduces token cost.

Large updates—such as reorganizing nested containers, merging State branches, generating code snippets, or performing multimodal analysis—occur offline, outside the LLM's token budget.

8.3 Categories of Actions

Actions may be categorized into several functional groups:

(1) State Manipulation Actions

- insert, update, delete values;
- recompute bindings;
- compress or prune history;
- externalize or rehydrate memory.

(2) Container Structural Actions

- create child containers;
- destroy obsolete containers;
- merge siblings;
- clone optimized container variants;
- wrap legacy objects.

(3) Multimodal Processing Actions

- generate embeddings;
- summarize images;
- extract audio features;
- process video frames;
- route content to specialized models.

(4) Tool and Code Actions

- invoke code callbacks;
- run system utilities;
- generate optimized code snippets;
- register new tool functions;
- override or intercept legacy methods.

(5) Communication Actions

- message sibling or parent containers;
- broadcast coordination signals;
- negotiate memory or task allocation;
- participate in emergent protocol formation.

(6) Scheduling and Control Actions

- request wake-ups;
- adjust heartbeat frequency;
- enter sleep mode;
- trigger conditional events.

These categories collectively support deterministic, auditable, and highly expressive control.

8.4 Dynamic Action Discovery and Expansion

Containers may autonomously:

- add new actions as they gain capabilities,
- generate domain-specific tools,
- prune obsolete actions,
- wrap external APIs into structured actions,
- inherit actions from parent containers,
- promote child actions to the parent interface.

This makes the action surface **evolutionary**, **extensible**, and **self-improving**.

8.5 Safe Execution and Sandboxing

Actions can be executed within a sandboxed environment, ensuring:

- deterministic state updates,
- isolation of risky behaviors,
- reversible or transactional state transitions,
- inspection of pre- and post-execution effects.

This ensures the system remains stable even as containers evolve.

8.6 Interaction with Token Budgets

Actions may incur token costs—especially those that expand State or increase context size. Containers may:

- decline costly actions when budget is low,
- choose more efficient alternatives,

- compress history before performing a large update,
- negotiate with sibling containers for shared token resources.

This introduces a functional relationship between **actions**, **token economics**, and **adaptive behavior**.

8.7 Actions as Cognitive Operators

In the architecture, Actions are not merely commands—they are **cognitive operators** used by the AI to:

- restructure memory,
- reorganize container graphs,
- delegate work,
- modify Purpose,
- shape their own cognitive trajectory.

This gives the model agency over its own structure, memory, and capabilities.

8.8 Summary of Action Module

The Action module provides the operational backbone of a context container. It enables deterministic state mutation, multimodal processing, distributed communication, structural self-modification, and adaptive scheduling—all through minimal-token inputs. Combined with Purpose and State, Actions make each container a self-governing unit capable of persistent cognition, evolution, collaboration, and large-scale autonomous operation.

9. Recursive Composition and Holonic Structure

The architecture supports **recursive composition**, enabling containers to be nested, wrapped, merged, split, cloned, pruned, or transformed to create arbitrarily deep and flexible hierarchies of cognitive units. This holonic structure is fundamental to the invention, allowing complex AI systems to emerge organically from simple, composable parts.

A recursively composable container (a "holon") functions simultaneously as:

- an autonomous cognitive unit,
- a component within a larger cognitive system,
- a host for subordinate cognitive units.

This dual identity mirrors biological and organizational systems and allows scalable, layered cognition.

9.1 Holon-as-Unit and Holon-as-System

Every container can serve as:

- a **standalone agent**, with its own Purpose, State, and Actions;
- a **sub-component** of a larger reasoning structure;
- a **container for further holons**, forming nested cognitive hierarchies.

This recursive symmetry allows:

- feature layering,
- specialization of sub-holons,
- multi-scale reasoning layers,
- compartmentalized memory systems,
- modular task pipelines.

9.2 Wrapping and Extension

A container can wrap:

- another container (decorator-like extension),
- a legacy code object or class,
- an external API or subsystem.

Wrapping allows a container to:

- enhance behavior,
- override or intercept methods,
- insert monitoring or quality checks,
- gradually replace or optimize underlying subsystems.

This enables the **AI-ization of existing codebases** without refactoring.

9.3 Merging, Splitting, and Cloning

Containers may undergo structural transformations:

- **Merge:** Combine sibling holons to unify or reconcile State.
- **Split:** Produce specialized derivative holons for sub-tasks.
- **Clone:** Create identical variants for swarm-balancing or redundancy.
- **Differentiation:** Allow clones to diverge based on Purpose or heuristics.

These operations allow modular evolution and rapid exploration of solution spaces.

9.4 Graphs of Holons

Holons may form:

- **trees**, supporting hierarchical control;
- **DAGs**, supporting layered computation;
- **meshes**, supporting distributed consensus;
- **swarm graphs**, supporting evolutionary optimization;
- **heterogeneous graphs**, where nodes use different models or modalities.

Edges represent:

- message flows,
- shared State references,
- dependency relationships,
- coordination pathways.

This graph-based architecture enables sophisticated multi-agent behavior.

9.5 Holonic Bundles for Model Invocation

Multiple holons may be bundled together and transmitted to a model as a **single composite workspace**. Bundling supports:

- multimodal fusion,
- joint reasoning across related units,
- token-efficient transmission of container groups,
- global structural updates in one pass.

Selective bundling allows the system to:

- include only relevant subgraphs,
- compress low-priority branches,
- externalize deep history,
- dynamically adjust context size.

9.6 Self-Referential and Meta-Holons

Holons may contain:

- evaluators for other holons,
- quality-control supervisors,
- schedulers,
- memory compaction modules,
- meta-optimizers.

Meta-holons can:

- observe sibling or child holons,
- prune or promote them based on performance,
- adjust their Purpose or token budgets,
- reconfigure container graphs.

This enables emergent governance structures.

9.7 Emergent Topology Evolution

Holon graphs may evolve based on:

- task difficulty,
- discovery of better strategies,

- token-budget constraints,
- performance feedback,
- container specialization.

The system may autonomously:

- grow new branches of holons,
- prune obsolete subtrees,
- restructure clusters,
- rebalance workloads across the swarm.

This allows the architecture to adapt continuously in an open-ended fashion.

9.8 Summary of Recursive Composition

Recursive composition transforms simple containers into complex, adaptive, and self-organizing AI systems. By allowing each holon to act as both unit and system—while enabling wrapping, nesting, merging, splitting, and autonomous evolution—the architecture provides the structural backbone for scalable, modular, and long-lived artificial cognition.

10. Serialization and Efficient Model Invocation

The architecture includes a specialized system for **serializing holons and holon graphs** into compact, predictable structures for transmission to AI models. Serialization transforms Purpose, State, and Action modules—along with any nested sub-containers—into token-efficient formats designed for high interpretability by large language models (LLMs) and multimodal systems.

This serialization pipeline is central to enabling long-context model utilization, multi-holon bundling, and scalable swarm reasoning.

10.1 Design Goals of Serialization

Serialization is engineered to:

- minimize token usage,
- produce strongly structured and predictable layouts,
- support merging or splitting container graphs,
- enable fast deserialization into running contexts,
- maintain schema-level consistency over time,
- support multimodal embeddings and cross-references,
- preserve action signatures and tool metadata,
- scale efficiently to hundreds or thousands of holons.

The result is a compact, deterministic representation of large cognitive states.

10.2 Standard JSON Serialization

Holons may be serialized into JSON using a standardized schema including:

- container identifiers,
- Purpose configuration objects,
- State structures,
- Action definitions and schemas,
- nested sub-holon references,
- communication channels,
- heartbeat and scheduling metadata,

- token budget data,
- historical windows.

JSON provides a clear, readable base format.

10.3 Token-Oriented Object Notation (TOON)

In some embodiments, holons are serialized using **Token-Oriented Object Notation (TOON)**—a compact, low-entropy external format optimized for transformers. TOON:

- reduces repeated tokens,
- collapses structural overhead,
- stores identifiers in compressed forms,
- normalizes common field names,
- packs values into predictable positions.

This significantly reduces token cost and improves computational efficiency.

10.4 Bundling Multiple Holons

Serialization supports bundling:

- entire holon subgraphs,
- swarms of cooperating holons,
- segmented pipelines,
- model-routed subsets.

A single call to an LLM may include dozens or hundreds of holons in a unified structure.

10.5 Schema Validation and Predictability

The system ensures predictable structure by enforcing:

- fixed field ordering,
- schema-validated keys,
- deterministic nesting rules,
- strict typing for Action parameters,
- clear separation of Purpose, State, and Action.

Predictability dramatically improves model attention efficiency.

10.6 State Delta Serialization

To minimize bandwidth, containers may:

- serialize only State deltas,
- elide redundant historical windows,
- compress embeddings,
- replace multimodal blocks with references.

This supports high-frequency heartbeat operation at scale.

10.7 Memory Externalization and Rehydration

Holons may offload parts of their State into external storage. When needed:

- externalized memory is rehydrated,
- embeddings are restored,
- sub-containers are pulled back into context.

This allows container graphs larger than the model's immediate context budget.

10.8 Cross-Reference Graph Encoding

Serialization includes:

- pointers to sibling holons,
- parent–child references,
- communication channel bindings,
- shared embedding pools,
- routing metadata.

Models can reason across holons as a unified graph.

10.9 Multimodal Component Encoding

State may include:

- embeddings for images, audio, code, or video,
- compressed metadata for external files,

- multimodal references for later retrieval.

The serialization layer ensures consistent formatting.

10.10 Attention Utilization Amplification Through Structured Serialization

Large-context LLMs (e.g., 1–2M tokens) often fail to fully use their window due to:

- unstructured inputs,
- verbose natural-language summaries,
- repeated text patterns,
- irregular layout,
- noisy logs.

By contrast, holon serialization produces:

- **highly predictable schemas**,
- **low-entropy token sequences**,
- **deterministic field ordering**,
- **explicit cross-references**,
- **minimal redundant content**,
- **tight packing of State**,
- **clear boundaries between cognitive units**,
- **strong relational cues for the attention mechanism**.

This dramatically increases the *effective* attention bandwidth. In many embodiments, systems may successfully bundle 20–50 holons, or large subgraphs, into a single invocation for a long-context model while maintaining high reasoning fidelity.

This enables true multi-agent parallel reasoning inside a single model call.

10.11 Serialization for Evolutionary Swarms

During swarm-balancing:

- variants may be serialized into a batch,
- evaluated in parallel by a single LLM invocation,
- scored together,
- pruned or propagated.

Serialization enables evolutionary loops at scale.

10.12 Summary of Serialization and Model Invocation

The serialization system provides a compact, deterministic, and token-efficient encoding of holons and holon graphs. Predictable structure amplifies model attention, supports large-scale bundling, enables multimodal workflows, and forms the backbone of efficient multi-agent reasoning across large-context AI models.

11. Model-Agnostic Interface Layer

The system includes a **model-agnostic interface layer** that abstracts away differences between AI backend models—language, vision, audio, video, code, embedding, hybrid, or future modalities. This layer allows containers to operate independently of any particular model architecture, API format, or vendor, enabling seamless switching, orchestration, and interoperability.

The interface layer ensures that a container's Purpose-driven routing instructions are fulfilled without requiring the container to know *how* a model is invoked or *what* model type is currently in use.

11.1 Goals of the Interface Layer

The interface layer is designed to:

- abstract model-specific APIs and formats,
- unify request/response handling,
- support multimodal routing,
- handle load balancing across models,
- enable dynamic or Purpose-driven model selection,
- maintain stability as underlying models evolve,
- reduce coupling between containers and backend systems.

It effectively provides **model polymorphism** for AI cognition.

11.2 Supported Model Types

The interface layer can route containers to any of the following:

- **Language models (LLMs)** for text reasoning, planning, or instructions.
- **Vision models** for images, video frames, OCR, or spatial understanding.
- **Audio models** for transcription, speech analysis, or feature extraction.
- **Video models** for temporal-spatial reasoning.
- **Embedding models** for high-dimensional representations.
- **Code models** for generation, refactoring, or static analysis.
- **Hybrid multimodal models** combining several capabilities.

- **Future model classes** unknown at the time of filing.

This allows heterogeneous systems to operate under a unified architecture.

11.3 Purpose-Driven Model Routing

Containers may specify:

- preferred models,
- fallback models,
- modal requirements (e.g., "requires vision"),
- cost-sensitive routing (e.g., "use cheaper model unless confidence too low"),
- specialized routes for sub-containers.

The interface layer resolves these preferences into operational routing decisions.

11.4 Request Normalization

Before a model is invoked, the interface layer transforms container bundles into model-specific formats:

- JSON or TOON structures are normalized,
- multimodal embeddings are injected or linked,
- irrelevant fields are pruned,
- system messages or scaffolding prompts are added as required.

Normalization ensures consistent results even across models with different API shapes.

11.5 Response Parsing and Validation

After model invocation, the interface layer:

- validates structured action outputs,
- ensures type correctness,
- checks schema conformance,
- rejects malformed or unsafe requests,
- normalizes responses into a uniform internal representation.

This ensures predictable behavior across all supported models.

11.6 Cross-Model Orchestration

The interface layer supports:

- routing different sub-containers to different models in parallel,
- aggregating multimodal outputs into a unified State update,
- delegating tasks to specialized models (e.g., vision → LLM → code),
- hybrid workflows where one model's output feeds another.

This enables multi-stage reasoning pipelines across heterogeneous systems.

11.7 Automatic Backend Migration

If a model is deprecated, scaled down, or updated, the interface layer may:

- transparently switch containers to a newer or equivalent model,
- preserve routing policies during migration,
- re-normalize serialization formats,
- test new model candidates before switching.

This ensures long-term system compatibility.

11.8 Failure Handling and Fallback Logic

In case of:

- model timeouts,
- malformed responses,
- low confidence signals,
- cost overruns,
- unavailable modalities,

the interface layer may:

- retry with the same model,
- route to a fallback model,
- reduce request complexity,
- invoke local actions instead,
- escalate to supervising containers.

This makes the system robust under variable conditions.

11.9 Extensibility

New model classes may be integrated by:

- defining a normalization schema,
- mapping response formats,
- establishing routing heuristics,
- registering model capabilities.

Containers automatically gain the ability to use new models without modification.

11.10 Summary of Model-Agnostic Interface Layer

The model-agnostic interface layer provides the abstraction necessary for containers to function independently of any specific AI model. By handling routing, normalization, parsing, fallback, and orchestration, it enables a flexible, resilient, and future-proof architecture that supports multimodal, multi-model AI cognition at scale.

12. Heartbeat-Driven Cognitive Cycle

The **heartbeat-driven cognitive cycle** is a central mechanism that shifts AI behavior from reactive, prompt-based interactions to proactive, chronological, and self-managed operation. Instead of only responding when prompted, containers may wake periodically, process their State, execute Actions, update memories, evaluate goals, or trigger structural changes.

Heartbeats create a predictable temporal substrate for long-running autonomous AI systems.

12.1 Purpose of the Heartbeat System

The heartbeat mechanism enables containers to:

- maintain long-lived tasks,
- update rolling memory windows,
- refine or compress State over time,
- schedule future work,
- perform periodic checks or maintenance,
- react to environmental changes,
- self-optimize based on token budgets,
- evolve internal structure via spawning or pruning holons.

This establishes a continuous cognitive process rather than isolated prompt-to-response cycles.

12.2 Heartbeat Scheduling

Each container may define, via its Purpose:

- fixed intervals (e.g., every 1 second),
- variable intervals (based on urgency or workload),
- predictive scheduling ("wake me when this condition is met"),
- budget-sensitive scheduling (reduced frequency during low token budgets),
- event-driven wake-ups (triggered by sibling or parent containers),
- cascading heartbeats (higher-level holons waking sub-holons).

The scheduler consolidates these into a global timeline.

12.3 Heartbeat Execution Cycle

On each heartbeat, the system:

1. Identifies containers marked for processing.
2. Runs State resolution (summaries, pruning, binding updates).
3. Serializes the relevant container graph (JSON or TOON).
4. Routes it to an appropriate model.
5. Parses model outputs.
6. Executes Actions to update State.
7. Re-assesses scheduling needs.

Each cycle represents a full cognitive step.

12.4 Adaptive Duty Cycling

To conserve resources, the heartbeat system may:

- temporarily lower frequency for inactive containers,
- increase frequency for urgent or high-priority tasks,
- batch multiple containers into a single invocation,
- skip updates when no changes are required,
- fully pause holons that enter long-term sleep.

Duty cycling transforms the architecture into an energy- and token-efficient cognitive graph.

12.5 State Evolution During Heartbeats

Heartbeats provide opportunities for containers to:

- refine memory summaries,
- reorganize or prune nested sub-containers,
- evaluate long-term progress,
- detect stale or redundant processes,
- initiate or terminate helper holons,
- migrate memory to more cost-efficient formats,
- generate code optimizations for future cycles.

This results in continuous self-improvement.

12.6 Heartbeats as a Multi-Container Coordination Mechanism

The scheduler can:

- align related containers onto shared intervals,
- coordinate synchronized updates across subgraphs,
- create cascading wake patterns for workflows,
- ensure upstream/downstream consistency,
- support swarm-balancing evolution phases.

This enables complex multi-agent behaviors that remain coherent over time.

12.7 Triggered and Conditional Heartbeats

Containers may request:

- conditional wake-ups ("wake me when container X changes its State"),
- deferred wake-ups ("wake me in 15 heartbeats"),
- event-driven wake-ups (external signals, API triggers),
- budget-triggered wake-ups (regain tokens before continuing work),
- consensus-triggered wake-ups (after swarm convergence).

This allows fine-grained control of container lifecycles.

12.8 Integration with Model Routing

Heartbeat-driven updates are tightly coupled with the model-agnostic interface layer:

- each heartbeat may route to different models depending on Purpose,
- containers may shift models over time,
- multimodal processing may occur in staggered cycles,
- heavy computations may be offloaded to code or external tools.

The heartbeat system acts as the orchestrator of heterogeneous model workflows.

12.9 Heartbeats in Evolutionary and Swarm Systems

During swarm-balancing phases, heartbeats regulate:

- parallel evaluation of cloned variants,

- mutation cycles,
- survival/pruning phases,
- propagation of optimal holons.

The heartbeat sequence effectively acts as the *clock* for evolutionary computation.

12.10 Summary of Heartbeat-Driven Cognitive Cycle

The heartbeat system transforms static or reactive AI behaviors into **continuous, stateful, self-managing cognition**. It enables long-horizon reasoning, autonomous evolution, multi-agent coordination, adaptive memory management, and efficient use of computation. Heartbeats provide the temporal foundation upon which the entire holonic architecture operates.

13. Tool Invocation and Code Callback Integration

The architecture provides a unified mechanism for invoking external tools, running code callbacks, integrating system utilities, and bridging between AI-driven behavior and conventional software logic. This system treats tools and code functions as **first-class cognitive operations**, available to any container through its Action module.

Unlike ad-hoc tool-use frameworks that depend on verbose natural language descriptions, the invention uses **minimal-token structured invocation**, ensuring efficient interaction even when tools trigger large or complex operations.

13.1 Goals of the Tool/Callback Layer

The design supports:

- efficient offloading of heavy computation,
- seamless integration with existing codebases,
- deterministic state mutation via callbacks,
- symbolic and typed action schemas,
- security, sandboxing, and rollback support,
- AI-generated tool definitions and dynamic expansion,
- code generation and self-improvement behaviors.

The result is a clean interface where AI and conventional code interact reliably.

13.2 Minimal-Token Tool Invocation

Instead of natural-language instructions, tools are invoked through compact objects:

```
{"action": "tool_name", "params": {...}}
```

Tools may:

- update State internally,
- return data to be incorporated on the next heartbeat,
- spawn or destroy containers,

- generate code snippets for future use.

This minimizes token expenditure while maximizing capability.

13.3 Typed Parameters and Schema Enforcement

Each tool defines:

- a parameter schema,
- expected types,
- required fields,
- optional arguments,
- return-value schemas.

The system automatically validates:

- type correctness,
- schema conformance,
- default value handling,
- safety constraints.

This prevents malformed or ambiguous tool requests.

13.4 Code Callback Execution

Tools may correspond directly to code callbacks in the host environment. These callbacks may:

- mutate container State,
- run external processes,
- access databases or APIs,
- generate optimized code for future invocation,
- perform computationally heavy tasks.

The callback environment is sandboxed to preserve determinism and safety.

13.5 AI-Generated Tools and Dynamic Registration

Containers may autonomously:

- generate new tool definitions,
- produce code snippets (Python, JS, etc.),
- test them in sandboxed execution,
- register them as new actions,
- deprecate or prune obsolete tools.

This enables **self-improving tool ecosystems**.

13.6 Legacy Code Integration and Method Redirection

Containers wrapping legacy classes may:

- intercept incoming method calls,
- redirect them through the Action system,
- merge return values into container State,
- override or replace methods with AI-optimized versions,
- maintain compatibility with existing APIs.

This enables incremental AI augmentation of traditional software without refactoring.

13.7 Tool Invocation Within Container Graphs

Tools may operate:

- on a single container,
- on sibling containers,
- across entire subgraphs,
- across the whole swarm during global updates.

Categories include:

- memory manipulation,
- multimodal processing,
- scheduling updates,
- container lifecycle operations,
- structural evolution actions.

13.8 Safety, Sandboxing, and Transactionality

Tool execution may occur in isolated environments with:

- transactional state changes,
- rollback protection,
- side-effect monitoring,
- capability restrictions.

If a tool misbehaves, the system can:

- revert State,
- throttle or disable the tool,
- escalate to supervisor containers.

13.9 Cross-Model and Cross-Modal Tool Orchestration

Tools may trigger workflows involving multiple models:

- embeddings → LLM → code model → vision model → LLM,
- hybrid multimodal transformation pipelines,
- model-specific post-processing.

The interface layer normalizes these operations into a single cohesive cycle.

13.10 Summary of Tool Invocation and Code Callback Integration

The tool and callback system provides a structured, extensible, and token-efficient bridge between AI reasoning and conventional computation. By supporting typed schemas, sandboxed execution, dynamic tool generation, and seamless legacy integration, it transforms containers into powerful hybrid cognitive units capable of invoking code, orchestrating tools, and evolving their own operational surface over time.

14. Inter-Container Communication Framework

The architecture includes a robust communication framework that enables containers to exchange information, coordinate actions, negotiate resources, and form higher-level collective behaviors. Communication occurs through structured, minimal-token messages that allow containers to interact without sharing full context windows.

This communication layer enables distributed cognition, swarm intelligence, and emergent protocol formation across holonic systems.

14.1 Goals of the Communication Framework

The system is designed to support:

- efficient cross-container messaging,
- coordination of parallel tasks,
- propagation of state changes,
- negotiation over shared resources (e.g., token budgets),
- collaborative problem-solving,
- multi-agent consensus-building,
- emergent communication protocols.

The framework treats communication as a first-class cognitive primitive.

14.2 Message Structure

Messages include:

- **sender ID**,
- **recipient ID(s)**,
- **schema-defined payload**,
- **timestamp or heartbeat index**,
- **optional priority and routing metadata**.

Payloads may be:

- symbolic commands,
- structured data objects,
- multimodal references,

- state deltas,
- coordination signals.

14.3 Communication Channels

Containers may use:

- **direct messaging** (one-to-one),
- **broadcast channels** (one-to-many),
- **topic-based channels** (publish/subscribe),
- **parent–child lanes** (hierarchical messaging),
- **swarm-wide consensus channels**.

Channels may be persistent or dynamically created.

14.4 Routing and Delivery

The communication layer supports:

- asynchronous delivery,
- synchronous coordination for joint tasks,
- priority-based routing,
- conditional routing (e.g., based on Purpose or State),
- inter-model delivery (e.g., LLM $\rightarrow$ vision $\rightarrow$ code model),
- cascading messages for multi-step workflows.

Routing is Purpose-driven and may evolve over time.

14.5 Communication-Driven State Updates

Messages can trigger:

- State mutations,
- activation of dormant holons,
- creation or destruction of sub-containers,
- role changes,
- updates in resource allocation,
- changes to scheduling or heartbeat frequency.

Thus, communication acts as a catalyst for structural evolution.

14.6 Negotiation and Token Economics

Messages may include:

- token-budget requests,
- memory-sharing proposals,
- competitive bidding for computational resources,
- cost estimates for upcoming updates.

Containers may:

- compete for resources,
- collaborate to minimize total system load,
- redistribute token budgets based on performance.

This introduces a lightweight economic layer.

14.7 Emergent Protocol Formation

Through repeated interactions, containers may develop:

- compressed symbolic communication codes,
- shorthand signals for common tasks,
- shared representations,
- non-human-readable internal languages.

The system does not assume fixed communication semantics—protocols may evolve.

14.8 Security and Verification

Communication is sandboxed and validated:

- schema validation of messages,
- authorization checks,
- message integrity guarantees,
- anti-loop and anti-flood protections.

Supervisory holons may audit communication patterns for anomalies.

14.9 Communication Across Heterogeneous Models

Messages may be routed to:

- LLM-driven containers,
- vision-based containers,
- code-generating containers,
- multimodal processors,
- external systems.

Cross-modal communication ensures coordinated workflows.

14.10 Summary of Communication Framework

The inter-container communication system enables distributed, coordinated, and emergent cognition across holonic networks. By providing minimal-token, structured messaging with routing, negotiation, and protocol-formation capabilities, the framework supports scalable multi-agent intelligence and self-organizing behavior.

15. Swarm-Balancing, Evolution, and Optimization

The architecture supports a powerful evolutionary mechanism—**swarm-balancing**—in which multiple holons with shared Purpose compete, cooperate, specialize, or fuse to produce optimized cognitive structures. This enables continuous self-improvement, adaptive task specialization, and emergent collective intelligence.

Swarm-balancing treats holons not as static agents but as evolving entities capable of replication, mutation, evaluation, and selection.

15.1 Goals of Swarm-Balancing

Swarm-balancing is designed to:

- explore diverse solution strategies in parallel,
- identify the most efficient or effective variants,
- prune suboptimal or redundant holons,
- propagate optimal configurations across the swarm,
- introduce diversity for robustness,
- support open-ended cognitive evolution.

This creates a self-optimizing ecosystem of cognitive units.

15.2 Parallel Instantiation of Variant Holons

A parent container may spawn multiple variants, each with:

- slight Purpose tweaks,
- alternative heuristics,
- different token-budget strategies,
- distinct model-routing preferences,
- varied history-window configurations,
- unique action surfaces.

These holons run **in parallel**, often on staggered heartbeats.

15.3 Independent State Evolution

Each variant develops its State independently:

- different memory summaries,
- divergent multimodal embeddings,
- unique internal representations,
- distinct structural transformations,
- variable optimization paths.

State divergence is expected—and desired—for exploratory breadth.

15.4 Evaluation Metrics and Performance Scoring

Holons may be scored using:

- task-specific accuracy or success metrics,
- token efficiency,
- model invocation cost,
- latency or speed,
- structural compactness,
- correctness validation by supervisor containers,
- inter-container consensus.

Multiple scoring strategies may be active simultaneously.

15.5 Selection and Pruning

Once evaluation is complete, the swarm may:

- prune low-performing variants,
- pause or recycle redundant holons,
- retain only high-performing branches,
- allow partial merging of promising but incomplete holons.

This reduces load on the system and consolidates gains.

15.6 Cloning and Propagation of Optimal Holons

When a variant outperforms others, the system may:

- clone it across the swarm,

- replace weaker holons with the optimized one,
- propagate its structure, Purpose, and State to sibling tasks,
- convert it into templates for future holon creation.

This spreads optimized cognition throughout the system.

15.7 Mutation and Diversity Maintenance

To prevent premature convergence, the system supports:

- purposeful mutation of heuristics,
- deliberate diversification of resource strategies,
- randomization in scheduling or routing,
- introduction of hybridized holons combining traits.

These techniques maintain exploratory capacity.

15.8 Cross-Holonic Cooperation

Variants may:

- share partial State summaries,
- exchange local discoveries,
- coordinate through communication channels,
- specialize into complementary roles.

This transforms competitive evolution into cooperative optimization.

15.9 Evolution Under Token Economics

Token budgets create natural selective pressures:

- efficient holons survive longer,
- memory-heavy holons prune or compress themselves,
- lean variants propagate more easily,
- budget constraints steer evolutionary paths.

This produces **computational natural selection**.

15.10 Swarm-Level Convergence

Over many heartbeats, the swarm may collectively reach:

- converged policies,
- optimized State structures,
- emergent protocols,
- stable holonic configurations.

Convergence triggers system-wide performance improvements.

15.11 Summary of Swarm-Balancing and Evolution

Swarm-balancing introduces evolutionary principles into the architecture. By spawning variants, evaluating performance, pruning failures, propagating successes, and maintaining diversity, the system continuously adapts, improves, and reorganizes itself. This transforms holonic networks into **living cognitive ecosystems** capable of sustained optimization and emergent intelligence.

16. Legacy Code Integration and System Augmentation

The architecture enables seamless **AI augmentation of existing software systems** without requiring refactoring, rewriting, or restructuring legacy codebases. This is achieved through dynamic wrapping, interception, State mirroring, and Purpose-driven transformation of existing classes, functions, APIs, and data structures into holons.

Legacy systems can be progressively enhanced, monitored, optimized, or replaced by holonic intelligence while maintaining API compatibility and operational continuity.

16.1 Goals of Legacy Integration

The system is designed to:

- introduce AI reasoning into existing codebases,
- maintain backward compatibility,
- enable gradual transformation rather than wholesale rewrites,
- provide real-time monitoring and optimization,
- automatically generate or refine code over time,
- convert software components into self-improving cognitive units.

This extends the architecture to real-world production systems.

16.2 Dynamic Wrapping of Legacy Classes

Any conventional class or object may be wrapped in a holon via:

- tag-based annotations,
- dynamic proxies,
- metaclass hooks,
- decorator-style wrappers,
- runtime interception.

The holon becomes the authoritative interface while delegating to the legacy object as needed.

16.3 Method Interception and Redirection

Wrapped objects provide a gateway for AI-driven control:

- incoming method calls flow through the holon's Action layer,
- the holon may modify parameters,
- inspect or adjust State,
- override behavior based on Purpose or heuristics,
- log the invocation for later analysis,
- decide whether to call the legacy function or a replacement.

This establishes a dynamic bidirectional link between code and cognition.

16.4 State Mirroring and Coherence Management

Legacy object state may be:

- mirrored into the holon's State module,
- tracked for drift between code and AI representations,
- reconciled during heartbeats,
- transformed into more efficient or symbolic formats.

Holons may detect inconsistencies and:

- correct them,
- escalate to supervisor holons,
- replace failing components,
- adjust their own internal representations.

16.5 Gradual AI-Driven Replacement of Legacy Logic

Over time, holons may:

- generate optimized code snippets replacing legacy behavior,
- benchmark legacy vs. AI-generated outputs,
- test new logic in sandbox mode,
- phase out inefficient code paths,
- elevate optimized logic into primary tool functions.

This supports a **progressive migration** toward AI-driven code.

16.6 Hybrid Execution and Delegation Control

Holons can dynamically decide whether to:

- fully delegate to native code,
- execute AI-generated code instead,
- merge outputs from both for validation,
- switch strategies based on token or compute budgets,
- fall back to legacy execution on failure.

This creates resilient and adaptive hybrid software.

16.7 External API Wrapping

Holons can also wrap external APIs, enabling:

- normalization of inconsistent endpoints,
- AI-driven retries or error handling,
- structured responses merged into State,
- multimodal interpretation of returned data,
- dynamic generation of API-specific tools.

This turns external services into holon-compatible components.

16.8 Monitoring, Audit, and Telemetry

Legacy-integrated holons may:

- track method frequencies,
- measure latency and cost,
- detect anomalies,
- log important events,
- generate performance audits,
- refine heuristics based on telemetry.

This transforms legacy software into an observable and optimizable substrate.

16.9 Safety and Containment Measures

Holons interacting with legacy systems operate within safeguards:

- transactional state updates,
- rollback protection,
- sandbox execution of generated code,
- capability restrictions,
- escalation pathways to supervisor holons.

This ensures reliability in production environments.

16.10 Summary of Legacy Integration

The architecture enables AI to progressively wrap, monitor, optimize, and eventually transform existing software systems. By treating legacy components as holons (or hosts for holons), the system provides a non-invasive path to AI augmentation, hybrid execution, dynamic replacement, and long-term evolution toward intelligent, self-improving code ecosystems.

17. Resource Management and Token Economics

The architecture incorporates a structured system of **resource management** in which computational capacity—primarily represented through token budgets—is treated as a quantifiable, allocatable, and optimizable resource. Containers may reason about their own resource usage, make trade-offs, negotiate with other containers, and adjust their behavior based on economic incentives encoded in their Purpose.

This creates a foundation for emergent efficiency strategies and self-regulating cognitive ecosystems.

17.1 Token Budgets as Computational Currency

Each container may be assigned a **token budget**, which may include:

- **Maximum Context Size** – the upper bound of serialized State.
- **Current Footprint** – the present size of the container's data.
- **Per-Heartbeat Allocation** – a recurring allowance consumed when updating.
- **Reserve Capacity** – optional emergency headroom.

Token budgets act as a form of **computational currency** that containers must manage wisely.

17.2 Budget-Aware Cognitive Behavior

Containers may examine their budgets and respond by:

- pruning or summarizing historical windows,
- externalizing deep memory to long-term storage,
- compressing multimodal embeddings,
- restructuring nested sub-containers,
- delaying non-essential work,
- reducing heartbeat frequency,
- delegating tasks to sibling holons.

This introduces a form of **self-interested efficiency**, encouraging lean, cost-effective cognition.

17.3 Inter-Container Resource Negotiation

Containers may negotiate or compete for resources through structured communication. Examples include:

- requesting additional token allocation from parent holons,
- proposing memory-sharing deals with siblings,
- bidding for resources when performing expensive operations,
- reducing their own footprint to accommodate global constraints.

This supports **distributed resource governance**.

17.4 Dynamic Budget Adjustment

Supervisory holons may adjust budgets based on:

- task importance,
- container performance,
- evolutionary fitness scores,
- system-wide load levels,
- operational policy.

Holons may grow, shrink, or stabilize their resource envelopes over time.

17.5 Cross-Holonic Token Flows

Token budgets may be transferred between holons as:

- incentives for good performance,
- penalties for inefficiency,
- collateral for requested operations,
- rewards for emergent optimization.

This creates **token flows** analogous to economic exchange within a digital ecosystem.

17.6 Budget-Constrained Evolution and Swarm Behavior

During swarm-balancing phases:

- efficient variants may survive longer,

- large-footprint variants may be pruned early,
- lean holons may propagate more easily,
- high-cost mutations may require explicit justification.

The system naturally favors holons that optimize both performance and resource expenditure.

17.7 Multi-Level Resource Arbitration

Resource arbitration may occur at multiple levels:

- **Local** (within a single holon),
- **Sibling-Level** (cooperative or competitive),
- **Parent–Child** (hierarchical governance),
- **Swarm-Level** (global optimization),
- **System-Level** (infrastructure constraints).

This allows resource management to adapt to different organizational structures.

17.8 Externalized Resource Markets (Optional Embodiment)

In certain embodiments, token budgets may be managed by:

- centralized economic controllers,
- decentralized markets between holons,
- dynamic auctions for scarce resources,
- credit or debt systems for deferred work.

These mechanisms allow complex emergent economies to form.

17.9 Resource Modeling for Non-AI Agents

The architecture does not require an AI backend to participate in resource dynamics. A container may:

- wrap a human operator,
- interface with a rule-based engine,
- coordinate with deterministic software systems.

In such cases, humans or external systems act as the "intelligence" within the holon, still subject to token budgets, resource constraints, and economic incentives.

17.10 Summary of Resource Management and Token Economics

The architecture transforms computational capacity into a manipulable, allocatable, and optimizable resource. Through token budgets, negotiation, dynamic adaptation, and evolutionary pressure, containers learn to manage their own memory, optimize their behaviors, and contribute to a stable, efficient, and self-regulating cognitive ecosystem.

18. Security, Safety, Isolation, and Failure Modes

The architecture includes a comprehensive security and safety framework designed to ensure predictable, trustworthy, and resilient operation of holons in both isolated and large-scale deployments. The framework addresses vulnerabilities in model outputs, code execution, memory mutation, inter-container communication, and interaction with external or legacy systems.

Holons operate under strong guarantees of isolation, transactional integrity, and controlled interaction, enabling reliable autonomous and semi-autonomous behavior.

18.1 Goals of the Security and Safety Framework

The system is designed to:

- contain faults within individual holons,
- prevent unsafe model outputs from causing uncontrolled behavior,
- enforce least-privilege access to tools and State,
- guarantee transactional and reversible updates,
- detect anomalies and corruption,
- maintain trust boundaries between containers,
- ensure graceful degradation under failure.

These measures protect system integrity across all operational layers.

18.2 Holon Isolation and Capability Boundaries

Each holon operates within strict capability boundaries:

- restricted access to tools,
- limited State mutation rights,
- sandboxed execution environments for code callbacks,
- explicit permission structures for inter-holon communication,
- constrained ability to spawn or destroy other holons.

This prevents runaway or unintended behavior.

18.3 Transactional State Updates

Every State mutation may occur within a transaction that provides:

- atomicity,
- rollback on failure,
- audit logs,
- verification checks before commit,
- delta-based serialization for efficient rollback.

If an update fails, the holon reverts to its last known-good State.

18.4 Sandboxed Tool Invocation

Tools and code callbacks execute in isolated environments with:

- restricted file system access,
- CPU/memory limits,
- network sandboxing (if applicable),
- deterministic side-effect tracking,
- execution timeouts.

Malicious or malfunctioning tools cannot corrupt container graphs.

18.5 Verification of Model Output

Model output is validated through:

- schema enforcement for Action structures,
- type checking for parameters,
- policy validation against Purpose constraints,
- safety filters (e.g., preventing destructive operations without authorization),
- cross-holon consistency checks.

Malformed or unsafe outputs are rejected automatically.

18.6 Anomaly Detection and Holon Quarantine

Holons may be quarantined if they exhibit:

- abnormal resource consumption,
- inconsistent State patterns,
- repeated schema violations,
- suspicious inter-holon communication,
- error cascades or runaway spawning behavior.

Quarantined holons may:

- be paused,
- isolated from siblings,
- placed under supervision,
- inspected or rolled back,
- terminated if unfixable.

18.7 Failure Modes and Graceful Degradation

Potential failures include:

- corrupted State,
- deadlocks in communication,
- tool execution failures,
- malformed model output,
- external API breakdowns,
- heartbeat scheduling failures.

The system responds with:

- fallbacks to previous States,
- reduced-frequency heartbeats,
- switching to alternative tools or models,
- rehydrating externalized memory,
- escalating to supervisory holons.

18.8 Safety Policies Embedded in Purpose

Purpose modules may include safety constraints such as:

- restricted actions,
- limits on spawn/destroy permissions,

- budget ceilings,
- approval requirements for high-risk operations,
- escalation rules for sensitive behaviors.

This embeds safety directly into the cognitive goals of the holon.

18.9 Secure Communication and Anti-Abuse Protections

Inter-holon communication includes:

- signature-based validation,
- anti-flooding and rate limits,
- loop-prevention mechanisms,
- access controls on broadcast channels,
- integrity checks for data payloads.

This prevents malicious or runaway messaging.

18.10 Legacy System Safety and Containment

When interacting with legacy systems, the architecture ensures:

- safe redirection of method calls,
- validation of external return values,
- controlled parameter mutation,
- fallback paths to original code,
- isolation of AI-generated logic from production-critical systems.

This allows safe progressive augmentation without jeopardizing existing functionality.

18.11 Human-in-the-Loop Safety Options

In some embodiments:

- humans may act as holons,
- human approval may be required for high-risk actions,
- human-reviewed summaries may validate State transitions,
- holons may escalate unclear decisions to human supervisors.

This adds external oversight when necessary.

18.12 Summary of Security, Safety, and Isolation

The system provides multilayered protections through isolation boundaries, sandboxing, transactional integrity, schema validation, anomaly detection, and controlled interaction with external systems. These mechanisms ensure that holons operate safely, predictably, and resiliently, even under failure or adversarial conditions.

19. External System Integration and Distributed Deployment

The holonic architecture supports integration with external systems, distributed computing environments, multi-node execution topologies, and hybrid on-prem/cloud infrastructures. Holons may operate on a single device, across a network, or within large-scale distributed systems while maintaining consistent State, Purpose, and Action semantics.

This section describes how holons interoperate with diverse computational substrates and coordinate across distributed environments.

19.1 Goals of External Integration and Distributed Deployment

The system is designed to:

- seamlessly connect holons to external processes and data sources,
- distribute holons across multiple compute nodes,
- support parallelization and load balancing,
- preserve consistency across network boundaries,
- enable latency-aware routing and scheduling,
- integrate with cloud, edge, and hybrid systems,
- coordinate execution across heterogeneous models or hardware.

These capabilities allow the architecture to scale horizontally and vertically.

19.2 Distributed Holon Graphs

Holon graphs may be partitioned and deployed across multiple nodes. Each partition maintains:

- local State integrity,
- subgraph-specific communication channels,
- routing metadata for cross-node messages,
- distributed heartbeat scheduling.

Nodes may operate independently while cooperating through structured inter-node communication.

19.3 Network-Aware Heartbeat Scheduling

Heartbeat frequencies may adjust dynamically based on:

- network latency,
- node availability,
- local resource constraints,
- operational load,
- historical performance.

Holons may request slower or faster heartbeats depending on their dependency on remote State.

19.4 Remote State Access and Synchronization

Holons may access remote State through:

- request/response protocols,
- cached summaries,
- delta synchronization,
- streaming State channels,
- snapshot restoration.

Synchronization may be:

- eventual,
- strongly consistent (with locking or versioning),
- conflict-resolved through heuristics or supervised rules.

19.5 Distributed Memory and Embedding Pools

State and embeddings may be sharded across:

- local memory,
- distributed caches,
- object stores,
- vector databases,
- long-term archive storage.

Holons rehydrate data on demand when needed by Action or Purpose.

19.6 Multi-Model Distributed Execution

The architecture supports routing holons or sub-tasks to:

- specialized LLMs,
- vision models,
- audio or speech systems,
- embedding models,
- code generation models.

Model selection may be:

- Purpose-driven,
- cost-sensitive,
- load-balanced,
- based on historical performance.

19.7 Hybrid Cloud/Edge Deployments

Holons may operate across:

- cloud servers,
- edge devices,
- on-prem clusters,
- mobile devices.

Edge holons may:

- perform local inference,
- compress State before transmitting upstream,
- manage sensory or real-time data.

Cloud holons may:

- handle coordination,
- store historical memory,
- perform heavy multimodal operations.

19.8 Integration with External Data Sources

Holons may interface with:

- databases,
- APIs,
- file systems,
- event streams,
- telemetry services,
- IoT sensors.

Integration uses structured Actions and tool callbacks with schema validation.

19.9 Distributed Swarm Balancing

Swarm-balancing may occur across nodes:

- variants may be spawned on different machines,
- evaluations may be run locally or centrally,
- high-performing variants may migrate or replicate across nodes,
- low-performing variants may be terminated.

This supports global optimization in distributed systems.

19.10 Fault Tolerance and Node Recovery

The system supports:

- node failure detection,
- holon migration,
- State restoration from checkpoints,
- degraded-mode operation until recovery,
- distributed consensus for critical decisions.

Holon graphs may automatically heal after partial network outages.

19.11 Secure Cross-Node Communication

Cross-node messages include:

- authentication tokens,
- encrypted payloads,

- routing metadata,
- integrity checks,
- anti-replay protections.

This allows safe operation across untrusted or semi-trusted network environments.

19.12 Summary of External Integration and Distributed Deployment

The architecture enables holons to operate seamlessly in distributed environments, coordinate across nodes, interact with external systems, and manage complex workflows across heterogeneous compute substrates. This supports large-scale deployment of holonic intelligence across modern cloud, edge, and hybrid infrastructures.

20. Example Implementations and Operational Workflows

This section provides concrete example implementations demonstrating how holons function in real systems. These examples serve as **enablement evidence**, illustrating how a practitioner can construct holons, deploy them, operate them under heartbeat scheduling, wrap legacy systems, and leverage swarm evolution.

The examples span chatbots, codebases, distributed agents, multimodal workflows, and large-scale parallel reasoning.

20.1 Example: Chatbot Wrapped in a Holon

Scenario: A standard chatbot is augmented by wrapping it in a holon.

Implementation Steps:

1. Initialize a holon with a Purpose describing conversation goals, tone, memory rules, and safety constraints.
2. Populate State with:
   - recent message history,
   - user profile (if allowed),
   - conversation context tree,
   - embeddings for long-term associations.

3. Define Actions such as:
   - `send_message` (transmits a reply),
   - `update_memory`,
   - `summarize_context`,
   - `schedule_followup`.

4. On each message (or heartbeat), serialize the holon and transmit to the model.
5. Model outputs one or more Actions instead of free-form text.
6. System executes Actions, updates State, and requests the next heartbeat.

Outcome:

- Conversation remains structured.
- Token usage stays minimal.

- Memory improves autonomously.
- The bot becomes a self-maintaining interactive agent.

20.2 Example: A Class or Object Wrapped in a Holon

Scenario: A legacy software class (e.g., `InvoiceProcessor`) is wrapped inside a holon.

Implementation:

1. The class is dynamically intercepted.
2. Method calls route into the holon's Action system.
3. State mirrors attributes: pending invoices, validation results, timestamps.
4. Purpose encodes rules like:
   - validation heuristics,
   - code generation triggers,
   - quality thresholds.
5. Holon monitors outputs and may replace inefficient methods with AI-generated tools.

Outcome:

- Legacy code gains AI reasoning without refactoring.
- The holon gradually improves or replaces methods.
- Execution remains safe through sandboxing.

20.3 Example: Multi-Holonic Pipeline for Data Processing

Scenario: A data ingestion pipeline is composed of holons.

Holons include:

- **Parser Holon** – parses raw data into structured formats.
- **Validator Holon** – checks schema consistency.
- **Transformer Holon** – performs feature extraction.
- **Enricher Holon** – fetches external metadata.
- **Publisher Holon** – writes results to storage.

Each holon:

- has a dedicated Purpose,
- manages independent State,
- requests tools for code execution,
- communicates with siblings for workflow coordination.

Outcome:

- The pipeline becomes adaptive.
- It self-corrects and evolves.
- It maintains itself under load.

20.4 Example: Heartbeat-Based Agent with No External Triggers

Scenario: A holon runs continuously without user interaction.

Implementation:

- Purpose includes periodic tasks: monitoring, scoring, indexing.
- Holon requests heartbeats every 10 seconds.
- On each heartbeat:
  - load updates are checked,
  - external data sources are polled,
  - previous State is analyzed,
  - tasks are executed or deferred.

Outcome:

- The system shifts from prompt-driven to time-driven.
- Holon becomes a continuously operating autonomous worker.

20.5 Example: 50-Holon Bundle for 2M-Token LLM Reasoning

Scenario: A large-context model processes a bundle of 50 holons simultaneously.

Serialization uses:

- strict schema ordering,
- TOON for compactness,
- cross-reference tables.

Each holon:

- annotates interdependencies,
- organizes subgraphs,
- exposes Actions and summaries.

Model processes all 50 simultaneously, producing coordinated outputs.

Outcome:

- True multi-agent parallel reasoning in one invocation.
- Effective attention spans dramatically improve.
- Eliminates the need for sequential calls.

20.6 Example: Swarm Evolution Cycle

Scenario: A holon evolves into optimized variants.

Steps:

1. Spawn 10 variants with modified Purpose heuristics.
2. Allow each variant to run for 100 heartbeats.
3. Score variants based on:
    - accuracy,
    - token efficiency,
    - error rates,
    - runtime.
4. Select top 2 variants.
5. Clone winners.
6. Replace losers.
7. Repeat.

Outcome:

- The system continuously evolves.
- Performance increases over time.
- Weak variants self-eliminate.

20.7 Example: Human Operator Acting as a Holon

Scenario: A human plays the role of a holon's "intelligence".

A UI panel provides:

- incoming structured State,
- Purpose goals and constraints,
- action slots the human can click,
- token budget indicators.

Human chooses Actions which the system executes.

Outcome:

- Humans can join holon networks.
- They obey the same token budgets and scheduling.
- Blended AI-human systems emerge.

20.8 Example: System-Wide Distributed Deployment

Scenario: Holons run across multiple machines.

Nodes host:

- localized holon groups,
- distributed memory shards,
- specialized routing functions.

Holons coordinate tasks like:

- multimodal embedding generation,
- cross-node data processing,
- distributed inference.

Outcome:

- Holons become a distributed operating system.
- They scale across heterogenous compute substrates.

20.9 Summary of Example Implementations

These implementations demonstrate real-world uses of holons in conversational agents, legacy integration, distributed systems, multi-agent reasoning, autonomous operation, swarm

evolution, and human-in-the-loop workflows. They collectively illustrate how the architecture can be applied directly, reliably, and flexibly in diverse environments.

21. Use Cases Across Domains

The holonic architecture is domain-agnostic and can be applied across a wide spectrum of computational, organizational, creative, and real-world systems. Because each holon is a structured cognitive container with clear Purpose, State, and Action surfaces, the architecture adapts naturally to any environment where reasoning, coordination, stateful computation, or autonomous workflows are required.

This section outlines representative use cases. These broaden the patent scope by demonstrating applicability across industries and technical disciplines.

21.1 Enterprise and Productivity Agents

Holons can manage:

- document drafting,
- email automation,
- calendar management,
- data entry and ETL,
- meeting summarization,
- workflow orchestration.

Multiple holons represent separate tasks, users, or departments, coordinating through structured communication.

21.2 Software Engineering and DevOps Automation

Holons may:

- analyze repositories,
- review code,
- propose patches,
- run CI/CD pipelines,
- monitor deployments,
- handle incident response,
- generate tests,
- evolve into specialized engineering assistants.

Legacy code can be progressively wrapped and upgraded over time.

21.3 Autonomous Research and Analysis Systems

Research-oriented holons may:

- read papers,
- extract findings,
- cross-analyze datasets,
- maintain long-term research memory,
- collaboratively summarize literature,
- spawn variants to explore competing hypotheses.

Swarm-balancing enables self-improving research loops.

21.4 Creative Tools and Content Generation

Holons may operate as:

- story planners,
- character managers,
- animation coordinators,
- composers,
- graphic design pipelines,
- video editing assistants.

Multiple holons collaborate across story arcs, timelines, or scene elements.

21.5 Robotics and Real-World Agents

Holons may represent:

- hardware modules,
- sensor suites,
- motion planners,
- object-recognition systems,
- high-level strategic control.

The structured Action layer allows safe tool invocation for real-world operations.

21.6 IoT, Edge Computing, and Embedded Systems

Distributed holons may:

- coordinate sensor networks,
- manage home automation,
- monitor industrial hardware,
- run predictive maintenance routines.

Edge holons operate with limited compute and transmit compressed updates upstream.

21.7 Financial Analysis and Trading Systems

Holons may:

- execute algorithms,
- track portfolios,
- evaluate risk,
- coordinate multi-model analysis,
- perform scenario simulations,
- optimize execution strategies.

Token budgets introduce computational cost-awareness.

21.8 Healthcare and Clinical Decision Support

Holons may:

- maintain longitudinal patient records,
- integrate multimodal medical data,
- run diagnostic reasoning,
- schedule follow-ups,
- coordinate between specialists.

Safety and validation features support sensitive workflows.

21.9 Education and Personalized Learning Systems

Holons may represent:

- student profiles,
- learning paths,
- skill maps,
- tutoring agents,
- content-generating assistants.

Swarm evolution enables personalized optimization.

21.10 Large-Scale Simulations and Game Worlds

Holons can represent:

- NPCs,
- factions,
- physics objects,
- dynamic story controllers,
- economy simulators.

Holons evolve individually or collectively, forming adaptive simulated societies.

21.11 Knowledge Management and Organizational Memory

Holons may:

- store corporate memory,
- maintain knowledge graphs,
- summarize documentation,
- coordinate across teams,
- preserve institutional knowledge over years.

Long-lived holons support multi-year operational continuity.

21.12 Human-in-the-Loop and Hybrid Intelligence Systems

Humans can act as holons or supervise holons. Use cases include:

- complex decision review,
- legal or policy evaluation,
- architecture oversight,

- creative direction.

Humans can be selectively inserted where needed.

21.13 Autonomous Codebases and Self-Maintaining Software

Entire software systems become:

- self-repairing,
- self-optimizing,
- self-monitoring,
- self-evolving.

Holons wrap, instrument, and eventually replace legacy logic incrementally.

21.14 Secure Multi-Tenant AI Platforms

Holons create natural tenant boundaries. Use cases:

- cloud AI hosting services,
- SaaS platforms with multiple agent groups,
- enterprise multi-division deployments.

Token budgets and Purpose policies enforce strong isolation.

21.15 Summary of Domain Use Cases

The holonic architecture applies seamlessly across conversational AI, devops automation, robotics, distributed systems, simulations, healthcare, finance, education, creative production, and hybrid human-AI workflows. This breadth demonstrates the flexibility and universality of the invention, supporting broad patent protection and enabling its adoption across nearly any domain.

22. Advantages Over Prior Art

The holonic architecture provides capabilities and behaviors that do not appear in existing agent frameworks, context-management systems, or LLM orchestration patterns. Unlike prior art—which typically relies on loosely structured prompts, static memory stores, or linear tool-use pipelines—this invention introduces a deeply modular, recursive, economic, evolutionary, and distributed cognitive substrate.

This section articulates the specific ways in which the invention surpasses prior approaches and provides non-obvious improvements that materially advance the state of the art.

22.1 Structured Cognitive Units vs. Unstructured Prompting

Prior Art:

- LangChain, AutoGPT, ReAct, BabyAGI, and similar frameworks inject large amounts of unstructured or semi-structured text into LLM prompts.
- Context is rewritten on every step, incurring high token cost.
- Models must infer structure from natural-language descriptions.

Advantage: The holonic architecture:

- enforces strict schemas,
- separates Purpose, State, and Action,
- eliminates unnecessary natural-language context,
- creates predictable, low-entropy structures.

This dramatically reduces model confusion and improves consistency.

22.2 Recursive Composability vs. Flat Agent Structures

Prior Art:

- Most agent frameworks treat agents as flat or loosely hierarchical entities.
- They rarely support recursive nesting or compositional self-expansion.

Advantage: Holons:

- can contain other holons,
- wrap existing holons,

- aggregate into arbitrarily deep graphs,
- support bottom-up and top-down cognitive flows.

This provides emergent hierarchical reasoning capabilities.

22.3 Action-Based Output vs. Free-Form Text

Prior Art:

- LLMs typically respond in natural language.
- Tool calls are often text-encoded.
- Systems rely on brittle parsing.

Advantage: Holons:

- require models to output structured Actions,
- use typed parameters,
- avoid expensive and error-prone NLP parsing,
- keep model output minimal and deterministic.

This reduces token cost and increases reliability dramatically.

22.4 Predictable Serialization vs. Context Chaos

Prior Art:

- Large context windows become cluttered with logs, summaries, and repeated text.
- LLM attention becomes diffuse.
- Long-context transformers underperform their advertised lengths.

Advantage: Holon serialization provides:

- deterministic field ordering,
- compact low-entropy layouts,
- TOON-optimized structures,
- explicit referencing between holons.

This increases *effective* attention bandwidth and allows meaningful reasoning across hundreds of thousands or millions of tokens.

22.5 Token Economics vs. Unlimited Growth

Prior Art:

- Memory expansion is largely uncontrolled.
- Systems often balloon in size without pruning.

Advantage: Holons:

- manage their own token budgets,
- prune or compress their memory autonomously,
- negotiate resources with siblings,
- align incentives with computational efficiency.

This introduces a novel economic substrate for self-regulation.

22.6 Evolution and Swarm Optimization vs. Static Agents

Prior Art:

- Agents usually operate in isolation.
- No built-in evolutionary mechanisms exist.

Advantage: Holons:

- spawn variants,
- evolve heuristics,
- prune failing instances,
- clone successful ones,
- converge toward optimized cognitive configurations.

This turns agent systems into adaptive, self-improving ecosystems.

22.7 Distributed Operation vs. Single-Node Execution

Prior Art:

- Many agent frameworks assume a single-machine environment.
- Cross-node communication is ad hoc.

Advantage: Holons:

- support distributed deployment natively,
- shard State and embeddings across nodes,
- coordinate through structured inter-node messaging,
- heal after partial network failures.

This enables scalable multi-node cognitive systems.

22.8 Legacy Code Wrapping vs. Full Rewrite Requirements

Prior Art:

- Integrating agents with legacy systems often requires major rewrites.
- LLMs cannot safely replace internal logic.

Advantage: Holons:

- wrap classes,
- redirect method calls,
- mirror object state,
- gradually replace code paths with AI-generated logic.

This provides incremental modernization without disruption.

22.9 Emergent Protocol Formation vs. Fixed Prompts

Prior Art:

- Systems rely on predefined instructions.
- No mechanism exists for emergent agent languages.

Advantage: Holons:

- can develop compressed communication codes,
- form symbolic or sublingual messages,
- coordinate using non-human-readable protocols,
- evolve communication strategies.

This yields more efficient collaboration at scale.

22.10 Deterministic Cognitive Surfaces vs. Implicit Reasoning

Prior Art:

- Reasoning is opaque.
- System behavior depends heavily on prompt phrasing.

Advantage: Holons:

- expose clear Purpose goals,
- surface State explicitly,
- define Action schemas,
- allow reproducible model invocation.

Behavior becomes inspectable and debuggable.

22.11 Full-Agent Bundling vs. Sequential Reasoning

Prior Art:

- Agents are often run sequentially, creating bottlenecks.

Advantage: Holons can be:

- bundled by the dozens,
- transmitted in one serialization block,
- reasoned about simultaneously.

This enables true multi-agent cognition in a single model invocation.

22.12 Distributed Human-AI Hybrids vs. Purely Automated Systems

Prior Art:

- Human oversight is usually outside the agent architecture.

Advantage: Holons:

- naturally incorporate humans as participants,
- enforce the same Purpose/State/Action structure,
- enable hybrid cognitive networks.

This expands the scope of the architecture beyond machine-only agents.

22.13 Summary of Advantages Over Prior Art

The holonic architecture provides structural, economic, evolutionary, and distributed capabilities unavailable in conventional agent frameworks. Its modular design, predictable serialization, token budgeting, swarm evolution, legacy code integration, and multi-model routing collectively represent a significant advancement, enabling robust, scalable, and adaptive artificial intelligence systems that exceed the limitations of prior art.

23. Variations, Embodiments, and Alternative Implementations

This section broadens the scope of the invention by enumerating alternative architectures, computational models, physical realizations, and hybrid systems that may implement the holonic principles described herein. These embodiments ensure that the invention is not limited to any particular hardware configuration, model type, serialization format, or software substrate.

All variations described below remain within the scope of the invention so long as they preserve the core principles of Purpose–State–Action segregation, structured serialization, and holon-level autonomy.

23.1 Non-AI Implementations

Holons do not require AI models to operate. In alternative embodiments, the cognitive decision-making role may be filled by:

- rule-based systems,
- finite-state machines,
- traditional software heuristics,
- symbolic logic engines,
- Bayesian or statistical models,
- constraint solvers,
- simple lookup tables.

These systems may still utilize Purpose, State, and Action architecture and participate in holon networks.

23.2 Human-Driven Embodiments

In some embodiments, a human acts as the intelligence within a holon. A UI panel provides:

- structured State summaries,
- Purpose constraints,
- selectable Actions with schema-defined parameters,
- token budget indicators.

The human selects actions, which the system executes while maintaining holonic semantics. Hybrid human–AI networks may emerge.

23.3 Hardware and Embedded Implementations

Holons may run on:

- microcontrollers,
- embedded processors,
- FPGAs,
- ASICs,
- robotic platforms,
- sensor hubs.

In these embodiments:

- State may be stored in non-volatile memory,
- Actions may invoke hardware-level operations,
- Purpose may encode hardware constraints,
- serialization may occur over low-bandwidth networks.

23.4 Distributed, Decentralized, and Federated Embodiments

Holons may operate in:

- peer-to-peer networks,
- decentralized clusters,
- federated learning environments,
- blockchain-backed systems,
- partially trusted or adversarial networks.

Consensus algorithms or supervisory holons may ensure global consistency.

23.5 Alternative Serialization Formats

Instead of JSON or TOON, holons may use:

- Protocol Buffers,
- MessagePack,

- CBOR,
- Cap'n Proto,
- custom binary formats,
- domain-specific compression algorithms.

Any deterministic, structured format may be substituted.

23.6 Alternative Purpose Models

Purpose may be defined using:

- natural-language constraints,
- symbolic rule sets,
- mathematical optimization targets,
- reinforcement-learning reward signals,
- hybrid behavioral specifications.

Purpose may include dynamic rewrites or self-generated updates.

23.7 Alternative State Models

State may include:

- symbolic graphs,
- dense embeddings,
- sparse vectors,
- structured relational tables,
- time-series buffers,
- multimodal latent spaces.

Different State layouts may be optimized for specific tasks.

23.8 Alternative Action Models

Action modules may:

- output natural language (if desired),
- produce bytecode or assembly commands,
- directly operate hardware via interrupts,

- invoke APIs or cloud services,
- generate code at runtime,
- perform symbolic manipulations.

Actions may also be learned or modified autonomously.

23.9 Variations in Holon Lifecycle Management

Holons may be:

- ephemeral,
- long-lived,
- spawned in large batches,
- pruned aggressively,
- paused for extended periods,
- rehydrated from checkpoints.

Holon death may be triggered by:

- resource exhaustion,
- evolutionary selection,
- policy violations,
- manual intervention.

23.10 Alternative Scheduling Models

Instead of heartbeat-based operation, holons may operate using:

- event-driven scheduling,
- rate-limited execution cycles,
- interrupt-based triggers,
- batch-processing windows,
- probabilistic or randomized intervals.

Scheduling models may also be hybridized.

23.11 Embodiments Without Recursive Nesting

In some use cases, holons may:

- avoid nested sub-holons entirely,
- operate only as flat sibling sets,
- use communication instead of nesting for coordination.

This may reduce structural complexity for simpler systems.

23.12 Embodiments With Deep Recursive Structures

Holons may instead use:

- arbitrarily deep hierarchies,
- dynamically growing and shrinking nested layers,
- self-wrapping behavior,
- fractal-like structural compositions.

These embodiments support extremely complex cognitive architectures.

23.13 Alternative Economic Models

Instead of token budgets, holons may use:

- CPU cycles,
- GPU time,
- memory quotas,
- network bandwidth,
- energy consumption,
- hardware wear coefficients.

Resource negotiation and self-governance remain intact.

23.14 Alternative Evolution and Selection Models

Swarm evolution may use:

- hill-climbing heuristics,
- genetic algorithms,
- gradient-based self-modification,
- Bayesian optimization models,
- novelty search,

- Monte Carlo tree search.

Selection pressures may vary widely between embodiments.

23.15 Holons as Modules in Traditional Software

Holons may replace:

- classes,
- objects,
- services,
- microservices,
- user sessions,
- state machines.

Traditional software may be incrementally converted into holonic systems.

23.16 Holons in Purely Physical Systems

Holons may represent:

- physical robots,
- assembly-line components,
- logistics hubs,
- smart vehicles,
- autonomous drones.

Purpose–State–Action becomes a universal control schema.

23.17 Hybrid Reality Embodiments

Holons may also operate in mixed environments where:

- some State is physical,
- some State is virtual,
- some Actions affect simulated worlds,
- some Actions affect real-world actuators.

This enables synthetic–physical hybrid ecosystems.

23.18 Summary of Variations and Embodiments

The holonic architecture supports a vast number of alternative implementations across hardware, software, hybrid systems, symbolic systems, human-in-the-loop designs, decentralized deployments, and multi-model intelligence stacks. All embodiments that preserve structured Purpose–State–Action semantics and holon-level autonomy fall within the scope of this invention.

24. Glossary and Definitions

This section provides clear definitions of key terms used throughout the specification. These definitions ensure terminological precision, eliminate ambiguity, and strengthen the legal clarity of the invention. All terms are intended to be interpreted broadly unless explicitly limited.

24.1 Holon

A modular cognitive unit containing Purpose, State, and Action modules. Holons may operate autonomously, communicate with other holons, evolve through swarm-balancing, or wrap external systems.

24.2 Purpose

A structured object describing the behavioral goals, heuristics, constraints, and optimization targets of a holon. Purpose guides decision-making and prioritization.

24.3 State

The full internal data representation of a holon, including memory, embeddings, operational history, multimodal content, nested sub-holons, and resource metrics.

24.4 Action

A schema-defined operation that a holon may request. Actions are structured, typed, and deterministic operations executed by the system or code callbacks.

24.5 Heartbeat

A system-driven or holon-requested update cycle in which serialized holons are transmitted to a model or reviewed for scheduled tasks.

24.6 Serialization

The process of encoding holons into deterministic formats such as JSON, TOON, or binary schemas, enabling efficient model invocation and cross-node transmission.

24.7 TOON (Token-Oriented Object Notation)

A compact, low-entropy serialization format designed to minimize token usage and maximize transformer attention efficiency.

24.8 Holon Graph

A network of holons linked through parent–child relationships, cross-references, communication channels, or task dependencies.

24.9 Swarm

A collection of holons—often variants of one another—operating in parallel, cooperating, competing, or evolving to optimize a shared Purpose.

24.10 Swarm-Balancing

An evolutionary mechanism where multiple holons explore different strategies, are evaluated, and undergo selection, cloning, pruning, or mutation.

24.11 Token Budget

A computational resource constraint specifying the maximum serialized size, per-heartbeat allowance, or lineage allocation available to a holon.

24.12 Token Economics

A resource-governance model in which holons negotiate, allocate, or trade token budgets as incentives for efficient behavior.

24.13 Delta Serialization

A method of transmitting only the changes (deltas) in State rather than the full State object.

24.14 Memory Externalization

The offloading of long-term or bulk State data into external storage, to be rehydrated when needed.

24.15 Tool

Any external code function, API endpoint, model call, or utility invoked by holons through structured Actions.

24.16 Callback

A local executable routine triggered by an Action, which may mutate State, perform computation, or interact with external systems.

24.17 State Mirroring

The reflection of legacy object or system state into a holon's State module for monitoring or gradual replacement.

24.18 Emergent Protocol

A communication pattern or symbolic mini-language developed autonomously between holons over repeated interactions.

24.19 Holon Wrapping

The process of encapsulating an existing system, object, class, API, or human interface inside a holon.

24.20 Cross-Node Communication

Inter-holon messaging across distributed systems, often secured, authenticated, and routed through network-aware channels.

24.21 Hybrid Holon

A holon whose cognitive processes are shared between AI models and human operators.

24.22 Supervisor Holon

A holon with elevated privileges responsible for monitoring, scoring, allocating resources, or orchestrating swarm behavior.

24.23 Sandboxing

A constrained execution environment preventing unsafe or uncontrolled tool or code behaviors.

24.24 Rollback

The process of reverting to a previous valid State after an error or failed transaction.

24.25 Context Amplification

A phenomenon in which structured serialization increases the effective usable context of long-context transformers.

24.26 Summary of Glossary

These definitions provide the structured terminology necessary for interpreting the invention consistently across embodiments, implementations, and future expansions.

25. Mathematical and Computational Formalization

This section provides formal models, pseudocode structures, and simplified mathematical descriptions of the holonic architecture. These formalizations demonstrate mechanizability and reinforce that the invention is not merely conceptual but can be instantiated in deterministic computational systems.

25.1 Holon as a State Machine

A holon (H) may be defined as a tuple:

H = (P, S, A)

Where:

- **P** is the Purpose module, defining constraints, heuristics, and objectives.
- **S** is the State module, containing internal data and nested holons.
- **A** is the Action module, providing callable operations.

A holon transitions according to:

$$S_{t+1} = f(P, S_t, A, I_t)$$

Where:

- I_t is input at heartbeat t (or null for self-triggered heartbeats),
- **f** is a deterministic transition function.

25.2 Heartbeat Scheduling Function

The heartbeat scheduler may be defined as:

$$h_{t+1} = g(S_t, P)$$

Where **g** outputs:

- next heartbeat time,
- sleep durations,
- conditional triggers.

Holons may request their own future wake-ups.

25.3 Structured Serialization Model

Serialization σ(H) produces a structured object:

σ(H) = (P', S', A')

Where **P', S', A'** are schema-enforced encodings.

Deserialization satisfies:

σ⁻¹(σ(H)) = H

Ensuring consistency across distributed systems.

25.4 TOON Entropy Reduction

Let **T** be the token sequence of a serialized holon.

TOON aims to minimize token entropy **H(T)** by enforcing:

- predictable field ordering,
- compressed identifiers,
- structural homogenization.

This improves transformer attention.

25.5 Effective Attention Amplification

Let **L** be the raw context window and **E** the effective usable attention.

In unstructured prompt systems:

- **E << L**

Under holonic serialization:

- **E ≈ L**

This demonstrates attention amplification due to structured input.

25.6 Inter-Holonic Messaging Model

A message **m** may be defined as:

m = (id_src, id_dst, payload)

Delivered via routing function:

R(m, G) → H_dst

Where **G** is the holon graph.

25.7 Evolutionary Fitness Function

Each holon variant H_i may receive a fitness score:

$$F(H_i) = w_1 P_i + w_2 E_i + w_3 C_i + w_4 S_i$$

Where:

- P_i = performance metric,
- E_i = token efficiency,
- C_i = correctness,
- S_i = structural compactness.

Selection probability:

$$Pr(H_i \text{ survives}) \propto F(H_i)$$

25.8 Token Budget Optimization as a Utility Function

Holons maximize utility:

$$U = \alpha B - \beta C$$

Where:

- **B** = budget preserved,
- **C** = cost of State expansion.

Subject to constraint:

$$|S| \leq B\_max$$

25.9 Distributed Holon Graph Partitioning

The holon graph **G** may be partitioned across nodes **N**:

$$\bigcup G_i = G$$

Partitions respect:

- locality constraints,
- resource availability,
- communication limits.

25.10 Pseudocode: Holon Execution Cycle

```
loop heartbeat:
    serialized = serialize(holon)
    response = model.invoke(serialized)
    actions = parse_actions(response)

    for action in actions:
        if validate(action):
            apply_action(action, holon)

    holon.state = update_state(holon)
    schedule_next_heartbeat(holon)
end loop
```

25.11 Pseudocode: Swarm Evolution Cycle

```
variants = spawn_variants(base_holon)

for v in variants:
    run_for_heartbeats(v, k)
    v.score = evaluate(v)

survivors = select_top(variants)
```

```
clones = clone(survivors)
replace_population(clones)
```

25.12 Summary of Formalization

This section provides mathematical, computational, and procedural models demonstrating that holons can be implemented as deterministic state machines, schedulable processes, distributed graph nodes, and evolutionary agents. These formalizations reinforce the reproducibility and technical completeness of the invention.

26. Reference Implementations

This section documents working implementations of the holonic architecture, demonstrating reduction to practice and providing concrete evidence of enablement. All implementations are publicly available and predate this filing.

26.1 HolonAI Library

Repository: https://github.com/NullCoward/HolonAI

A Python library providing the core Holon abstraction with full Purpose–State–Action architecture.

Key Features Implemented:

- **Purpose Module** (§6): Configurable purpose definitions via `.add_purpose()`
- **State Module** (§7): Dynamic self-concept via `.add_self()` with callable bindings
- **Action Module** (§8): Schema-derived actions with automatic metadata extraction
- **Serialization** (§10): JSON and TOON (Token-Oriented Object Notation) output
- **Token Management** (§17): Built-in token counting via tiktoken with budget enforcement
- **Recursive Composition** (§9): Nested holons supported natively

Code Example:

```python
holon = (
    Holon(name="TaskManager")
    .with_token_limit(4000, model="gpt-4o")
    .add_purpose("You are a task management assistant")
    .add_self({"user": "alice"}, key="context")
            .add_action(create_task,   name="create_task",
purpose="Create a new task")
```

```
    )
    prompt = serialize_for_ai(holon)
```

26.2 AI Chat Room

Repository: https://github.com/NullCoward/AIChatRoom

A multi-agent chat application where AI agents communicate autonomously using heartbeat-driven polling.

Key Features Implemented:

- **Heartbeat-Driven Cognitive Cycle** (§12): Agents polled on configurable intervals (1-10 seconds)
- **Inter-Container Communication** (§14): Agents send structured messages to each other
- **Token Budgeting** (§17): 10k total budget with distribution by attention percentage
- **Self-Concept** (§7): Flexible JSON store for agent-managed knowledge
- **Command System** (§8): Structured action outputs (speak, join_room, update_self_concept)
- **Swarm Behavior** (§15): Agents can create and spawn other agents

Architecture:

- Agent = Room paradigm (each agent owns its communication space)
- HUD (Heads-Up Display) system for context window management
- OpenAI Responses API integration with model routing

26.3 WikiHolon

Repository: https://github.com/NullCoward/WikiHolon

A self-maintaining AI-powered wiki system demonstrating autonomous content management.

Key Features Implemented:

- **Tool Invocation** (§13): Database operations via structured callbacks
- **Legacy Integration** (§16): SQLite persistence with auto-migrations
- **External System Integration** (§19): FastAPI REST endpoints

- **Autonomous Operation** (§12): Self-maintaining content updates

---

26.4 Summary of Reference Implementations

These implementations collectively demonstrate:

1. **Technical Feasibility** – The architecture can be instantiated in working software.
2. **Multiple Embodiments** – The same principles apply across different domains (library, chat, wiki).
3. **Reduction to Practice** – Actual, functioning code predates this filing.
4. **Enablement** – A person of ordinary skill can reproduce the invention using these references.

All source code is available under open-source licenses and may be examined to verify the claims made in this specification.